

Lecturer: Andrew Parkes
<http://www.cs.nott.ac.uk/~ajp/>

G52ADS 2014-15

Shortest Path & Dijkstra

Shortest path

- Given a graph with weights/distances on the edges
- Find the shortest route between two vertices u and v .
- It turns out that we can just as well compute shortest routes to ALL vertices reachable from u (including v).
 - This is called *single-source shortest path problem* for weighted graphs, and u is the source.

Dijkstra's Algorithm

- An algorithm for solving the single-source shortest path problem.
- Assume that weights are non-negative (though possibly zero)
- Think of the weights as distances, and the length of the path is the sum of the lengths of edges.

Remarks

- Will develop the algorithm by repeated refinement – will see the same example 3 times, each time, with more detail.
- To understand this algorithm (and others), after each stage, ask yourself
 - “what is now known to be true and that was not known in previous stages?”
 - such understanding forms the basis of “appropriate assertions and loop invariants” and is assessable

Example in “code perspective”

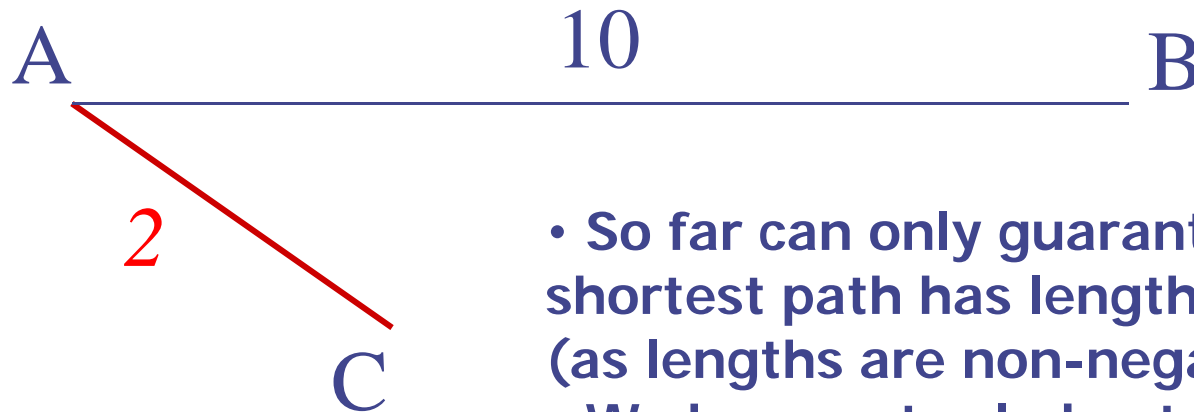
- Looking for shortest path from A to B
- Start from node A & find neighbours

A

Example

- Looking for shortest path from A to B
- So shortest path is 10


NO !!



- So far can only guarantee that the shortest path has length at least 2 (as lengths are non-negative)
- We have not ruled out the possibility of a path length $L(A,B)$ with $2 \leq L(A,B) < 10$


Example

- Which node should we expand next?
- Expand C as trying to rule out shortest paths
- Now know: shortest path A-C is 2




Example

- Next: expand D as it is
 - not yet expanded
 - the one with the shortest path and we are trying to rule out that $L(A,B)$ is in the range $[4 : 10]$



Example

- Now we have reached B with $L(A,B) = 6$
- Are we finished?
- Yes, in this case, as all nodes are expanded (except B itself)



Core Ideas

- The previous simple example contains the core ideas of Dijkstra
 - “expand” means “add neighbours to a working list”
 - expand nodes with the shortest known current path as this is the only node for which we know the distance is really the shortest possible
 - **do not prematurely assume that have found the shortest path to a node**

Dijkstra's algorithm

To find the shortest paths (distances) from the start vertex s :

- keep a priority queue PQ of vertices to be processed
- for each u in the PQ maintain $\text{dist}(s,u)$ as the shortest current known path length from s to u
 - e.g. keep an array with current known shortest distances from s to every vertex (initially set to be infinity for all but s , and 0 for s)
- always order the queue so that the vertex with the shortest distance is at the front.
 - Exercise: ensure that you understand, and can explain, **why** this must be done.

Dijkstra's algorithm

Loop while there are vertices in the queue PQ:

- dequeue a vertex u – from the front, “popMin”, hence with the least $\text{dist}(s,u)$
- expand node u :
 - recompute shortest distances for all vertices in the queue (i.e. not ‘closed’) as follows:
 - if there is an edge from u to a vertex v in PQ

$$\text{dist}(s,v) \leftarrow \min(\text{dist}(s,v) , \text{dist}(s,u) + w(u,v))$$

- close u , i.e. move to a “closed” list

Computing the shortest distance

If the shortest distance from s to u is $\text{distance}(s,u)$ and the length (weight) of the edge between u and v is $w(u,v)$, then the current shortest distance from s to v is $\text{distance}(s,u) + w(u,v)$.




Important

- Do **NOT** conclude have the shortest path to a node until it has moved to front of the PQ and been dequeued and moved to the closed list
- Now do the same example again but this time with the PQ done explicitly:


Example

- $PQ = \{A(0)\}$ Closed = $\{\}$
- Dequeue and expand A




Example

- $PQ = \{ C(2) , B(10) \}$ $Closed = \{ A(0) \}$
- Dequeue and expand C



Example


- $PQ = \{ D(4), B(10) \}$ $Closed = \{ A(0), C(2) \}$
- Dequeue and expand D & recompute B



$$= \min(10, 4+2)$$

Example

- $PQ = \{ B(6) \}$ Closed = $\{ A(0), C(2), D(4) \}$
- Dequeue and close B and conclude $L(A,B)=6$



Pseudocode for D's Algorithm

- PQ : priority queue of unvisited vertices prioritised by shortest recorded distance from source
- $PQ.reorder()$ reorders PQ if the values in *dist* change.

Pseudocode for D's Algorithm

```
PriorityQueue PQ = new PriorityQueue();
while (! PQ.isEmpty()){
    u = PQ.dequeue();
    if ( u == target ) return dist[u];
    for(each v adjacent to u){
        add v to the PQ if not present and not
        already closed, else update the distance using
        if(dist[v] > (dist[u]+weight(u,v)){
            dist[v] = (dist[u]+weight(u,v));
        }
    }
    add u to list of closed nodes
    PQ.reorder(); // because some distances changed
}
return INFINITY; // no path to target
```

Implementing the PQ

- Many choices:
- It is not quite a heap – as might need to access nodes other than the minimum in order to change the distance
- Might just live with duplicates – and check when remove nodes that they are not already closed
- See GoTa textbook, etc, for advanced options


Finding the Path

To make Dijkstra's algorithm to return the path itself, not just the distance:

- In addition to distances, maintain a "back pointer" $\text{back}(u)$ a pointer to the previous node in the best path to u
- By following the back pointers can rebuild the path
- In the beginning paths are empty
- When adding a expanding u gives a new node v then $\text{back}[v]=u$
- When re-assigning $\text{dist}(s,v)=\text{dist}(s,u)+\text{weight}(u,v)$ also re-assign $\text{back}(v)=\text{back}(u)$.

Example


- $PQ = \{A(0,-)\}$ Closed = $\{\}$
- Dequeue and expand A



the back pointer
from C to A


Example

- $PQ = \{ C(2,A) , B(10,A) \}$ $Closed = \{ A(0,-) \}$
- Dequeue and expand C



Example


- $PQ = \{ D(4,C) , B(10,A) \}$ $Closed = \{ A(0,-), C(2,A) \}$
- Dequeue and expand D & recompute B & back(B)



new back pointer

Example

- $PQ = \{ B(6,D) \}$
Closed = $\{ A(0,-), C(2,A), D(4,C) \}$
- Close B and optimal back path is D,C,A



Optimality of Dijkstra's algorithm

- So, why is Dijkstra's algorithm optimal (gives the shortest path)?
- Let us first see where it could go wrong.

What the algorithm does

- For every vertex in the priority queue, we keep updating the current distance downwards, until we remove the vertex from the queue.
- After that the shortest distance for the vertex is set.
- What if a shorter path can be discovered later?

Optimality proof

- Base case: the shortest distance to the start node is set correctly (0)
- Inductive step: assume that the shortest distances are set correctly for the first n vertices removed from the queue. Show that it will also be set correctly for the $n+1$ st vertex.

Optimality proof

- Assume that the $n+1$ st vertex is u . It is at the front of the priority queue and its current known shortest distance is $\text{dist}(s,u)$. We need to show that there is no path in the graph from s to u with the length smaller than $\text{dist}(s,u)$.
- Proof: by contradiction – but non-essential so moved to “appendix” for self-study if desired

Complexity

- Assume that the priority queue is implemented as a heap;
- At each step (dequeueing a vertex u and recomputing distances) we do $O(|E_u| * \log(|V|))$ work, where E_u is the set of edges with source u .
- We do this for every vertex, so total complexity is

$$O((|V| + |E|) * \log(|V|))$$

- Really similar to BFS and DFS, but instead of choosing some successor, we re-order a priority queue at each step, hence the extra $\log(|V|)$ factor.

Exercise

- You are **highly** recommended to
 - create some small to medium graphs and work through this algorithm
 - repeat working examples until you understand it fully and can do it 'by hand' quickly and easily
 - Dijkstra is a classic algorithm, and the same ideas appear in many other algorithms

Minimum Expectations

- Know and understand definition of shortest path and Dijkstra's algorithm
- Be able to apply it, by hand, to small graphs
- Be able to argue or explain why it does give the shortest path
 - Formal proof not needed

Module Wrap-up

- General theme: Describing, reasoning about, and reducing the run-time of algorithms:
 - “Describing” - Big oh family
 - “Reasoning about” – counting operations, using heights of trees, etc
 - “Reducing” – using binary search, divide and conquer, binary trees, avoiding repeated work, etc
- Specifics:
 - “Linear structures”: Arrays, vectors, linked lists, stacks, queues, hashmaps
 - “Tree structures”: BST, heaps, pre- post- in-order traversals
 - “Graphs”: breadth & depth-first search, MST, Dijkstra

(All of these are core CS.)

END

"Appendix"

- Material that is
 - slightly more difficult
 - not strictly required
 - but that might still illuminate other aspects, or act as a test of understanding

Context: Greedy Algorithms

- Dijkstra's algorithm: has a stage “pick the vertex to which there is the shortest path currently known at the moment.”
- This is a “greedy” strategy
 - For Dijkstra's algorithm, this turned out to be globally optimal: a shorter path to the vertex can never be discovered.
 - However there are (many) problems for which greedy strategies which are not globally optimal:

Example:

Non-optimal greedy algorithm

- Problem: given a number of coins, count the change in as few coins as possible.
- Greedy strategy: start with the largest coin which is available; for the remaining change, again pick the largest coin; and so on.
- Example: Coins $\{5, 2, 2, 2, 2\}$. “Change”: 8
 - the largest coin ‘5’ is a ‘fatal mistake’ – not part of any desired solution
- “Exercise”: find a fast greedy algorithm for the general version of this problem, or show one does not exist.
 - NOTE: \$1m prize for solving this “exercise” ☺ – see G53COM
 - Note: not a real exercise for G52ADS! – but see “millennium prize problems”

Diameter of a graph

- http://en.wikipedia.org/wiki/Distance_%28graph_theory%29
- Diameter of a connected undirected graph is the length of the “longest shortest path”
 - the maximum over all pairs of nodes a, b
 - of the length of the shortest path between a and b

Optimality proof

- Proof by contradiction: assume there is such a (shorter) path
- That path contains a vertex v_1 to which the shortest distance is set (it may be that $v_1=s$) which has an edge to a vertex v_2 to which the distance is not set (maybe $v_2=u$)
 - Exercise: why must there be such a v_2 if the path is to be shorter?



Optimality proof

- So the vertices from s to v_1 have correct shortest distances (inductive hypothesis) and v_2 is still in the priority queue.



Optimality proof

- So $\text{dist}(s, v1)$ is indeed the shortest path from s to $v1$. Current distance to $v2$ is $\text{dist}(s, v2) = \text{dist}(s, v1) + \text{weight}(v1, v2)$



Optimality proof

- If $v2$ is still in the priority queue, then $\text{dist}(s, v1) + \text{weight}(v1, v2) \geq \text{dist}(s, u)$



Optimality proof

- But then the path going through $v1$ and $v2$ cannot be shorter than $\text{dist}(s,u)$.
QED

